

LWE and Algebraic Variants

CSE 598 Final Project

Jordan Jones

April 2026

Abstract

We survey the Learning With Errors (LWE) problem and its algebraic variants, including Ring-LWE, Module-LWE, Polynomial-LWE, and Order-LWE. We begin by reviewing the standard LWE problem and the foundational hardness results of Regev, which establish a quantum reduction from worst-case lattice problems to LWE. We then describe how Ring-LWE replaces vectors with elements of polynomial rings, yielding dramatically more efficient cryptographic constructions while preserving worst-case hardness guarantees via ideal lattice problems. Finally, we discuss the unifying framework of “Algebraic LWE Revisited,” which shows that Ring-LWE reduces to the proposed algebraic LWE variants over number fields, and discuss applications to NIST post-quantum standards and fully homomorphic encryption.

Contents

1	Introduction and Motivation	3
2	Learning With Errors	3
2.1	Problem Definition	3
2.2	Establishing Hardness: Regev’s Reduction	3
2.2.1	Proof Sketch: The Iterative Construction	4
2.2.2	Worst-Case to Average-Case Reduction	4
2.2.3	Decision to Search Reduction	4
2.3	LWE-Based Public Key Encryption	5
3	Ring-LWE	5
3.1	Algebraic Preliminaries	5
3.2	The Canonical (Minkowski) Embedding	6
3.3	Main Theorem (Ring-LWE Hardness)	7
3.4	Pseudorandomness of Ring-LWE	7
3.5	Proof Outline	8
3.5.1	Reduction 1: $\text{LWE}_{q,\Psi} \rightarrow \mathfrak{q}_i\text{-LWE}_{q,\Psi}$ (Recover Secret mod Prime Ideal)	8
3.5.2	Reduction 2: $\mathfrak{q}_i\text{-LWE}_{q,\Psi} \rightarrow \text{WDLWE}_{q,\Psi}^i$	8
3.5.3	Reduction 3: $\text{WDLWE}_{q,\Psi \leq \alpha}^i \rightarrow \text{DLWE}_{q,\Upsilon_a}^i$ (Worst-Case to Average-Case) . .	8
3.5.4	Reduction 4: $\text{DLWE}_{q,\Upsilon}^i \rightarrow \text{DLWE}_{q,\Upsilon}$ (Hybrid Argument)	8
3.6	Efficiency Gains	8

4	More Algebraic Variants	9
4.1	Module-LWE	9
4.2	Polynomial-LWE	9
4.3	Order-LWE	9
4.4	Algebraic LWE Revisited: A Unifying Framework	10
4.4.1	Generalized $\mathcal{L}$ -LWE	10
4.4.2	Main Results	10
5	Applications	10
5.1	NIST Post-Quantum Cryptography Standards	10
5.2	Fully Homomorphic Encryption	11
6	Open Problems	11
7	Things I Learned	11
8	Updates	11
9	Conclusion	12

1 Introduction and Motivation

The Learning With Errors (LWE) problem, introduced by Regev [1], has become one of the central hardness assumptions in lattice-based cryptography. Its appeal stems from a worst-case to average-case reduction: solving LWE on random instances is at least as hard as approximating certain lattice problems—specifically the Shortest Independent Vectors Problem (SIVP) and the decision Shortest Vector Problem (GapSVP)—in the worst case.

While LWE yields powerful cryptographic primitives (public-key encryption, fully homomorphic encryption, attribute-based encryption, etc.), the standard formulation suffers from large key sizes: the public key is of size $\tilde{O}(n^2)$, which can be rather large for typical parameters in practice. This motivates the study of *algebraic variants* of LWE, which exploit additional ring structure to achieve compact keys and efficient arithmetic, while (ideally) preserving worst-case hardness guarantees.

In this survey we cover three main topics:

- (i) The standard LWE problem and Regev’s hardness reduction (Section 2).
- (ii) Ring-LWE and its hardness based on ideal lattice problems (Section 3).
- (iii) Further algebraic variants (Module-LWE, Polynomial-LWE, Order-LWE) and the unifying framework of Algebraic LWE Revisited (Section 4).

We conclude with applications to NIST post-quantum standards and fully homomorphic encryption (Section 5).

2 Learning With Errors

2.1 Problem Definition

Definition 2.1 (LWE Distribution). *Let $n \geq 1$ be an integer, p be an integer modulus, and χ be an error distribution over $\mathbb{Z}_p$. Given a secret $\mathbf{s} \in \mathbb{Z}_p^n$ and $\mathbf{a}_i \leftarrow \mathbb{Z}_p^n$ sampled uniformly, LWE samples are of the form*

$$(\mathbf{a}_i, b_i = \langle \mathbf{s}, \mathbf{a}_i \rangle + e_i) \quad \text{for } e_i \leftarrow \chi.$$

We write $\text{LWE}_{p,\chi}$ for this problem, parametrized by the modulus p and error distribution χ .

There are two main formulations:

Search-LWE. Given samples $(\mathbf{A}, \mathbf{b}^t)$ where $\mathbf{b}^t = \mathbf{s}^t \mathbf{A} + \mathbf{e}^t$, with $\mathbf{A} \leftarrow \mathbb{Z}_p^{n \times m}$ and $\mathbf{e} \leftarrow \chi^m$, find the secret $\mathbf{s}$.

Decision-LWE. Distinguish an LWE instance $(\mathbf{A}, \mathbf{b})$ from a uniformly random sample $(\mathbf{A}, \mathbf{u})$ where $\mathbf{u} \leftarrow \mathbb{Z}_p^m$.

2.2 Establishing Hardness: Regev’s Reduction

The foundational hardness result for LWE is due to Regev [1].

Theorem 2.2 (Regev 2005). *For $\varepsilon = \varepsilon(n)$, $p = p(n) \in \mathbb{Z}$, and $\alpha = \alpha(n) \in (0, 1)$ such that $\alpha p > 2\sqrt{n}$, there is an efficient quantum reduction from $\text{DGS}_{\sqrt{n} \cdot \varepsilon(\mathcal{L})/\alpha}$ to $\text{LWE}_{p, \Psi_\alpha}$.*

The reduction uses the Discrete Gaussian Sampler (DGS) oracle to approximate:

- The Shortest Independent Vectors Problem (SIVP): output a set of n linearly independent lattice vectors of length at most $\gamma(n) \cdot \lambda_n(\mathcal{L})$.
- The decision version of the Shortest Vector Problem (GapSVP): given a lattice $\mathcal{L}$, decide whether $\lambda_1(\mathcal{L}) \leq d$ (YES instance) or $\lambda_1(\mathcal{L}) > \gamma(n) \cdot d$ (NO instance).

Specifically, the hardness of approximating SIVP and GapSVP to within $\tilde{O}(n^{1.5})$ in the *worst case* is tied to the hardness of LWE, establishing the crucial **worst-case to average-case** connection.

2.2.1 Proof Sketch: The Iterative Construction

Informally, Regev's reduction uses an iterative construction. Starting from a “wide” Gaussian over the lattice, it takes iterative steps towards narrower Gaussians over the lattice:

1. **Classical step.** Use an LWE oracle and samples from the DGS sampler of the current iteration to solve a bounded distance decoding problem over the dual lattice. This effectively allows decoding at a larger radius (by a factor of roughly αp), which is what enables the transition to a narrower Gaussian in the next iteration.
2. **Quantum step.** Starting from a uniform superposition on the lattice cosets $\sum_{x \in L^*} |x\rangle$, apply reversible operations and add in gaussian error relative to the new width to obtain the state

$$\sum_{x \in L^*, z \in \mathbb{R}^n} \exp(-\pi \|rz\|^2) |x, x + z\rangle.$$

By using the BDD oracle from part one, we can erase x from the register and obtain the state $\sum f_{1/r}(z) |z\rangle$, where $f_{1/r}$ is the Fourier transform of $D_{L,r}$. By applying the quantum Fourier transform, we obtain the state $\sum D_{L,r}(y) |y\rangle$, and by measuring, we sample from the discrete Gaussian over the lattice parametrized by r . (*Here $f_{1/r}$ is a continuous Gaussian whose Fourier transform yields the discrete Gaussian $D_{L,r}$; this step uses the Fourier duality of Gaussians to convert a smooth distribution into a lattice distribution.*)

Starting from n^c samples from $D_{L,r}$ and letting $r_i = r \cdot (\sqrt{n}/(\alpha p))^i$, we repeat steps (1) and (2) n^c times to get n^c samples from D_{L,r_i} . Continuing this process $\approx 3n$ times gives the results from the main theorem.

2.2.2 Worst-Case to Average-Case Reduction

LWE also admits a *worst-case to average-case reduction*. The proof idea is to randomly shift samples: for $t \in \mathbb{Z}_p^n$,

$$f_t(a, b) = (a, b + \langle a, t \rangle),$$

which transforms $A_{s,\chi}$ into $A_{s+t,\chi}$, making a random LWE instance as hard as any specific one.

2.2.3 Decision to Search Reduction

There is also a decision-to-search reduction. Informally, given a sample $(\mathbf{a}, b)$, we transform it to $(\mathbf{a} + (k, 0, \dots, 0), b + k \cdot \ell)$ for $k, \ell \in \mathbb{Z}_p$. A sample distinguisher can test whether $\ell = s_1$ (the first coordinate of $\mathbf{s}$). Since there are $p = \text{poly}(n)$ possibilities for each coordinate, this can be done efficiently on each coordinate to recover the full LWE secret. Assuming the Search version is hard, this result implies the LWE distribution itself is pseudorandom.

2.3 LWE-Based Public Key Encryption

LWE admits the following public-key cryptosystem:

- **Gen**(1^n): Sample $\mathbf{s} \leftarrow \chi^n$, $\mathbf{A} \leftarrow \mathbb{Z}_q^{n \times m}$, and $\mathbf{e} \leftarrow \chi^m$. Set

$$\mathbf{b}^t = \mathbf{s}^t \mathbf{A} + \mathbf{e}^t \in \mathbb{Z}_q^{1 \times m}.$$

The public key is

$$\text{pk} = \begin{pmatrix} \mathbf{A} \\ \mathbf{b}^t \end{pmatrix} \in \mathbb{Z}_q^{(n+1) \times m},$$

and the secret key is $\text{sk} = \mathbf{s}$.

- **Enc**($u \in \{0, 1\}$): Choose a random vector $\mathbf{x} \leftarrow \{0, 1\}^m$. Output

$$\mathbf{c} = \begin{pmatrix} \mathbf{A} \\ \mathbf{b}^t \end{pmatrix} \mathbf{x} + \begin{pmatrix} \mathbf{0} \\ u \lfloor q/2 \rfloor \end{pmatrix} \in \mathbb{Z}_q^{n+1}.$$

- **Dec**($\mathbf{c}$): Compute

$$d = \mathbf{c}^t \cdot \begin{pmatrix} -\mathbf{s} \\ 1 \end{pmatrix} \in \mathbb{Z}_q.$$

Output 0 if $d \in [-q/4, q/4] \bmod q$, and output 1 otherwise.

Remark 2.3. *The public key is of size $\tilde{O}(n^2)$. This is not very efficient for practical applications, which motivates the study of algebraic variants.*

3 Ring-LWE

The key idea of Ring-LWE is to *replace vectors with ring elements in specific rings*, gaining compactness and efficiency.

3.1 Algebraic Preliminaries

Definition 3.1 (Cyclotomic Polynomial). *The m -th cyclotomic polynomial, written $\Phi_m(x) \in \mathbb{Z}[x]$, has degree $n = \varphi(m)$ and its roots are the primitive m -th roots of unity:*

$$\{\zeta \in \mathbb{C}^\times \mid \zeta^m = 1, \text{ord}(\zeta) = m\}.$$

Explicitly,

$$\Phi_m(x) = \prod_{i \in \mathbb{Z}_m^*} (x - \omega_m^i),$$

where $\omega_m = e^{2\pi i/m}$.

One can show that $\Phi_m(x)$ is the minimal polynomial of any primitive m -th root of unity and is irreducible over $\mathbb{Z}[x]$.

The Ring R . Going forward, we work with the ring

$$R = \mathbb{Z}[x]/\langle \Phi_m(x) \rangle,$$

the ring of integer polynomials modulo a cyclotomic polynomial. Since $\Phi_m(x)$ is irreducible over $\mathbb{Z}[x]$, R is a domain. Elements can be represented by degree $< n$ polynomials over $\mathbb{Z}[x]$, via the basis $\{1, \bar{x}, \dots, \bar{x}^{n-1}\}$.

Remark 3.2. Since R is a domain, it embeds into its field of fractions $K = \mathbb{Q}(\zeta_m)$.

Definition 3.3 (Ring-LWE Distribution). Let $R_q = R/qR$. For secret $s \in R_q$, sample $a \leftarrow R_q$ uniformly and error $e \leftarrow \chi$. A Ring-LWE sample is

$$(a, b = a \cdot s + e \bmod qR).$$

Number Field. A number field is a finite field extension of the rational numbers obtained by adjoining a root of a monic polynomial over $\mathbb{Q}$. The field extension $\mathbb{Q} \hookrightarrow K \cong \mathbb{Q}(\zeta_m)$ is obtained by adjoining a primitive m -th root of unity. The field $\mathbb{Q}(\zeta_m)$ is the splitting field of $x^m - 1$ over $\mathbb{Q}$, and hence the extension is normal. Because characteristic 0 implies separability, $\mathbb{Q} \hookrightarrow \mathbb{Q}(\zeta_m)$ is a Galois extension.

It can be shown that

$$\text{Gal}(\mathbb{Q}(\zeta_m)/\mathbb{Q}) \cong (\mathbb{Z}/m\mathbb{Z})^*,$$

where the isomorphism sends

$$[\ell] \mapsto \sigma_\ell,$$

with σ_ℓ defined by

$$\sigma_\ell(\zeta_m) = \zeta_m^\ell$$

for ℓ such that $\gcd(\ell, m) = 1$. This map is injective, since an automorphism of $\mathbb{Q}(\zeta_m)$ is determined by the image of ζ_m . Moreover, by the Galois property,

$$|\text{Gal}(\mathbb{Q}(\zeta_m)/\mathbb{Q})| = \varphi(m) = |(\mathbb{Z}/m\mathbb{Z})^*|.$$

So the map is an isomorphism.

Ring of Integers. The *ring of integers* $\mathcal{O}_K$ of a number field K is the set of elements $\alpha \in K$ that are roots of some monic polynomial in $\mathbb{Z}[x]$. For our cyclotomic field $K = \mathbb{Q}(\zeta_m)$, we have $\mathcal{O}_K = \mathbb{Z}[\zeta_m]$, which is precisely our ring R . More generally, an *order* $\mathcal{O} \subseteq K$ is a subring of $\mathcal{O}_K$ that is also a full-rank $\mathbb{Z}$ -module; the ring of integers is the unique maximal order.

Remark 3.4. We pass to the number field K because it admits embeddings into $\mathbb{R}^n$, giving a geometric interpretation of R . This allows ideals to be viewed as lattices and enables Gaussian error distributions and worst-case lattice reductions, which are not available purely within R .

3.2 The Canonical (Minkowski) Embedding

To obtain a lattice from a number field in a way that respects algebraic structure, we choose an embedding $\sigma_i : K \rightarrow \mathbb{C}$ that fixes $\mathbb{Q} \subseteq K$. There are $n = \varphi(m)$ such embeddings:

$$\sigma_i(\omega_m) = \omega_m^i \quad \text{for } i \in \mathbb{Z}_m^*.$$

The *canonical (Minkowski) embedding* is defined as:

$$\sigma(x) = (\sigma_1(x), \dots, \sigma_n(x)) \in \mathbb{C}^n.$$

Each $\sigma_i(x)$ can be identified as an element of $\mathbb{R}^2$ via $z \in \mathbb{Q}(\zeta_m) \mapsto (\text{Re}(z), \text{Im}(z))$, so $\sigma(x) \in \mathbb{R}^{2n}$ via this embedding. Under this embedding, addition and multiplication are coordinate-wise, and automorphisms simply permute coordinates.

Definition 3.5 (Ideal Lattice). *For an ideal $I \subseteq \mathcal{R}$, the ideal lattice $\sigma(I)$ is the image of I under the canonical embedding.*

Remark 3.6. *The error distribution in Ring-LWE is a product distribution on independent Gaussians in this embedding.*

Remark 3.7. *The domain of the non-noisy products $s \cdot a_i$ in Ring-LWE is most naturally defined over the codifferent ideal $R^\vee$, the trace-dual of R , defined by*

$$R^\vee = \{x \in K : \text{Tr}(xR) \subseteq \mathbb{Z}\}.$$

Under the canonical embedding, this corresponds to the dual lattice. This choice ensures that the Gaussian error behaves correctly under the Fourier transform used in worst-case reductions; using coefficient-based noise without this structure can lead to bad instantiations [3].

3.3 Main Theorem (Ring-LWE Hardness)

The main theorem consists of two component results:

1. **Worst-case hardness of search.** A quantum reduction from approximate SVP (worst case) on ideal lattices in R to the search version of Ring-LWE. This follows the general outline of Regev’s iterative quantum reduction to DGS (Section 2.2).
2. **Search/decision equivalence.** A *classical* reduction showing that the R-LWE distribution is pseudorandom, i.e., search Ring-LWE reduces to average-case decisional Ring-DLWE.

Theorem 3.8 (Lyubashevsky–Peikert–Regev [2]). *If it is hard for polynomial-time quantum algorithms to approximate the search version of the Shortest Vector Problem (SVP) in the worst case on ideal lattices in R , to a fixed $\text{poly}(n)$ factor, then any $\text{poly}(n)$ number of samples drawn from the R-LWE distribution are pseudorandom to any polynomial-time (quantum) attacker.*

3.4 Pseudorandomness of Ring-LWE

The search/decision equivalence (component 2 above) is also established in [2], which shows that $A_{s,\psi}$ itself is pseudorandom when s, ψ are chosen at random (and kept secret) from appropriate error distributions. This is established via a chain of four reductions:

$$\text{LWE}_{q,\Psi} \xrightarrow{\text{Automorphisms}} \mathbf{q}_i\text{-LWE}_{q,\Psi} \xrightarrow{\text{Search/Decision}} \text{WDLWE}_{q,\Psi}^i \xrightarrow{\text{Worst/Avg}} \text{DLWE}_{q,\Upsilon}^i \xrightarrow{\text{Hybrid}} \text{DLWE}_{q,\Upsilon}$$

Two subtleties arise in the ring setting that do not appear in standard LWE:

Recovering the full secret. In the ring setting, a DLWE oracle might only let us deduce the value of the secret s relative to *one* ideal factor $\mathbf{q}_i$ of $\langle q \rangle$. To recover the entire secret, we “shuffle” the $\mathbf{q}_i$ factors using the field automorphisms of K : since the Galois group acts transitively on the prime ideals $\mathbf{q}_i$ above q (a standard fact for Galois extensions: the $\mathbf{q}_i$ are conjugate under $\text{Gal}(K/\mathbb{Q})$), we can recover s modulo all prime factors and reconstruct s via the Chinese Remainder Theorem.

Error distribution randomization. The first component (worst-case hardness of search) establishes hardness of $\text{LWE}_{q,\psi}$ for non-spherical error distributions $\psi \in \Psi$, which may not be individually invariant under field automorphisms. However, the entire *family* Ψ is invariant, and so the reduction must randomize the error distribution itself.

3.5 Proof Outline

We now describe the four reductions in more detail.

3.5.1 Reduction 1: $\text{LWE}_{q,\Psi} \rightarrow \mathbf{q}_i\text{-LWE}_{q,\Psi}$ (Recover Secret mod Prime Ideal)

For a sample $(a, b) \leftarrow A_{s,\psi}$, we transform it into the sample $(\tau_K(a), \tau_K(b))$ where τ_K is a field automorphism such that $\tau_K(\mathbf{q}_j) = \mathbf{q}_i$. Use the oracle for $\mathbf{q}_i\text{-LWE}$ along with field automorphisms to recover $s \bmod \mathbf{q}_i$ for each prime factor $\mathbf{q}_i$ of $\langle q \rangle$. Reconstruct the full secret s via the Chinese Remainder Theorem.

3.5.2 Reduction 2: $\mathbf{q}_i\text{-LWE}_{q,\Psi} \rightarrow \text{WDLWE}_{q,\Psi}^i$

Definition 3.9. Let $i^- :=$ the largest element in $\mathbb{Z}_m^*$ less than i . Define the distribution $A_{s,\psi}^{i-}$ as follows: choose $(a, b) \leftarrow A_{s,\psi}$ and output $(a, b + h/q)$ where $h \in R_q$ is uniformly random and independent mod $\mathbf{q}_j R$ for all $j \leq i$, and 0 mod the remaining $\mathbf{q}_j R$.

Definition 3.10. $\text{WDLWE}_{q,\Psi}^i$: Given access to $A_{s,\psi}^j$ for arbitrary $s \in R_q^\vee$, $\psi \in \Psi$, and $j \in \{i^-, i\}$, determine j .

The proof idea for recovering $s \bmod \mathbf{q}_i R$ is to try each possible value, modifying the samples according to whether the value matches the correct one (follows the distribution A^{i-}) or not (using the distribution A^i). A $\text{WDLWE}_{q,\Psi}^i$ oracle tells us which is which.

3.5.3 Reduction 3: $\text{WDLWE}_{q,\Psi_{\leq \alpha}}^i \rightarrow \text{DLWE}_{q,\Upsilon_\alpha}^i$ (Worst-Case to Average-Case)

$\text{WDLWE}_{q,\Psi}^i \rightarrow \text{DLWE}_{q,\Upsilon_\alpha}^i$ (distinguishing $A_{s,\psi}^i$ vs $A_{s,\psi}^{i-}$ for random s, ψ).

The proof idea is to randomly shift the problem by adding a secret offset s' and extra noise e' to simulate the average-case distribution, mapping between the two distributions. The oracle's distinguishing advantage is preserved in expectation over the random shift, so whether this advantage is non-negligible or negligible allows us to deduce the original distribution.

3.5.4 Reduction 4: $\text{DLWE}_{q,\Upsilon}^i \rightarrow \text{DLWE}_{q,\Upsilon}$ (Hybrid Argument)

Suppose (S, Ψ) is a pair for which an oracle distinguishes $A_{s,\psi}$ vs uniform with non-negligible advantage. For the sequence $A_{s,\psi}^0, \dots, A_{s,\psi}^{m-1}$, uniform, there must exist some $i \in \mathbb{Z}_m^*$ for which the oracle distinguishes $A_{s,\psi}^i$ vs $A_{s,\psi}^{i-}$ with non-negligible advantage.

3.6 Efficiency Gains

The key efficiency improvements of Ring-LWE over standard LWE are:

- Each noisy product $b = a \cdot s + e$ gives n simultaneously pseudorandom values over $\mathbb{Z}_q$ rather than a single scalar.
- Polynomial multiplication can be performed in $O(n \log n)$ scalar operations using the FFT/NTT.

- Each sample $(a, b) \in R_q \times R_q$ can replace n samples $(a_i, b_i) \in \mathbb{Z}_q^n \times \mathbb{Z}_q$, reducing the public key size by a factor of n , collapsing keys to $O(n \log q)$ bits.

Security. The best known algorithms for SVP approximation on ideal lattices to within polynomial factors require exponential quantum time. However, one can obtain $2^{\sqrt{n}}$ approximations in quantum polynomial time.

4 More Algebraic Variants

Since the original Ring-LWE paper, additional variants of “algebraic” LWE over commutative rings have arisen.

4.1 Module-LWE

Module-LWE generalizes Ring-LWE to free modules of arbitrary rank over the ring R_q . That is, samples are tuples $(\mathbf{a}, b) \in R_q^d \times R_q$ where $\mathbf{a}$ is a vector of ring elements. It serves as a tradeoff between the efficiency of Ring-LWE and the security confidence closer to standard LWE.

4.2 Polynomial-LWE

Polynomial-LWE works over the ring $R_q = \mathbb{Z}_q[x]/\langle f(x) \rangle$ with samples of the form

$$(a(x), b(x) = a(x) \cdot s(x) + e(x) \mod f(x)),$$

where $e(x)$ is “small” in the coefficient basis.

Remark 4.1. *Poly-LWE is effectively Ring-LWE viewed through the coefficient embedding rather than the canonical embedding. While Ring-LWE relies on the geometry of the number field to define noise, Poly-LWE defines “smallness” directly on the polynomial coefficients. For power-of-two cyclotomics $f(x) = x^n + 1$, these two views are computationally equivalent. In general, the error lives in the codifferent ideal $R^\vee$ (corresponding to the dual lattice under the canonical embedding), but for power-of-two cyclotomics, $R^\vee = \frac{1}{n}R$. This linear relationship allows the coefficient-based noise to be geometrically equivalent to the canonical noise (up to the scalar $1/n$).*

4.3 Order-LWE

Definition 4.2. Let K be a number field with ring of integers $\mathcal{O}_K$. An order $\mathcal{O} \subseteq K$ is a subring that is a full-rank $\mathbb{Z}$ -module; that is, its $\mathbb{Q}$ -span equals K .

For example, for the number field $\mathbb{Q}(\sqrt{5})$, the ring of integers is $\mathcal{O}_K = \mathbb{Z}\left[\frac{1+\sqrt{5}}{2}\right]$, and $\mathbb{Z}[\sqrt{5}]$ is a (non-maximal) order contained in $\mathcal{O}_K$. Order-LWE is LWE over the quotient

$$\mathcal{O}_q = \mathcal{O}/q\mathcal{O},$$

with samples $(a, b = a \cdot s + e \mod q\mathcal{O})$ where $a, s, e \in \mathcal{O}_q$.

Remark 4.3. *Ring-LWE can be seen as a special case of Order-LWE by letting $\mathcal{O} = \mathcal{O}_K$, the full ring of integers.*

Remark 4.4. *Reductions between these variants were historically difficult to interpret due to the complexity of parameters and large blowup of error distributions.*

4.4 Algebraic LWE Revisited: A Unifying Framework

The paper “Algebraic LWE Revisited” [4] unifies and simplifies these variants, giving a general framework that encompasses all proposed algebraic LWE variants over number fields, using the hardness of the original Ring-LWE problem as a foundation for hardness of all algebraic LWE problems.

4.4.1 Generalized $\mathcal{L}$ -LWE

All of these problems share the same structure: given a number field K , an order $\mathcal{O} \subseteq K$, and certain modules or ideals over $\mathcal{O}$, the secret s , public samples a , and products $s \cdot a$ live in quotient structures of the form $\mathcal{I}/q\mathcal{I}$ where $\mathcal{I}$ is an ideal and q is the modulus. The map $(s, a) \mapsto s \cdot a$ is a bilinear operation over these modules.

Definition 4.5 (Generalized LWE). *$\mathcal{L}$ -LWE is parametrized by any lattice $\mathcal{L}$ of a number field K (a free additive group whose $\mathbb{Q}$ -span equals K). Let R be a ring, A, S, T be R -modules, $\langle \cdot, \cdot \rangle : A \times S \rightarrow T$ an R -bilinear map, and χ an error distribution over T . Then $\mathcal{L}$ -LWE samples are of the form*

$$(a, b = \langle a, s \rangle + e) \quad \text{for } a \in A, s \in S, e \leftarrow \chi \text{ (“small” noise)}.$$

This framework recovers the prior algebraic LWE problems:

1. **Ring-LWE**: when $\mathcal{L} = \mathcal{O}_K$ is the ring of integers of K .
2. **Poly-LWE**: when $\mathcal{L} = \mathbb{Z}[\alpha]^\vee$ for some $\alpha \in \mathcal{O}_K$.
3. **Order-LWE**: when $\mathcal{L} = \mathcal{O}$ for some arbitrary order $\mathcal{O}$ of K .

4.4.2 Main Results

Theorem 4.6 (Algebraic LWE Revisited [4]).

1. *There is a reduction from $\mathcal{L}'$ -LWE to $\mathcal{L}$ -LWE for any sublattice $\mathcal{L} \subseteq \mathcal{L}'$ of K , provided that the index $|\mathcal{L}'/\mathcal{L}|$ is coprime to q . This reduction preserves error distributions. Therefore, by taking $\mathcal{L}' = \mathcal{O}_K$ (the full ring of integers, corresponding to Ring-LWE), the worst-case hardness theorems and hardness guarantees are inherited by Order-LWE (by taking $\mathcal{L} = \mathcal{O}$).*
2. *There is a reduction from $\mathcal{O}$ -LWE for a wide class of orders $\mathcal{O}'$ to a single $\mathcal{O}$ -problem of higher rank in K (i.e., Module-LWE), over a suborder $\mathcal{O} \subseteq \mathcal{O}'$, giving hardness foundations for Module-LWE.*

The key takeaway is that Ring-LWE reduces to Order-LWE and Module-LWE problems with essentially no error degradation, or distributions related by linear transformations. All algebraic LWE problems over number fields are instances of a single unifying framework, and Ring-LWE gives a clean hardness foundation for them.

5 Applications

5.1 NIST Post-Quantum Cryptography Standards

The algebraic LWE variants form the foundation for the NIST Post-Quantum Cryptography (PQC) standards. The National Institute of Standards and Technology (NIST) is a U.S. government agency that evaluates and standardizes cryptographic schemes for real-world deployment. Modern standards use **Module-LWE**, which inherits some of the efficiency of Ring-LWE as well as hardness results. The standardized NIST module-based schemes include:

- **ML-KEM** (formerly CRYSTALS-Kyber): a key encapsulation mechanism.
- **ML-DSA** (formerly CRYSTALS-Dilithium): a digital signature algorithm.

5.2 Fully Homomorphic Encryption

Ring-LWE is used in modern Fully Homomorphic Encryption (FHE) schemes, enabling computation on encrypted data. The ring structure provides the algebraic operations needed for efficient homomorphic multiplication, and the compact representation keeps ciphertext sizes manageable.

6 Open Problems

1. **Classical reductions.** Can the original lattice reduction for LWE be made *classical* (i.e., without quantum computation)? Regev’s reduction is inherently quantum; removing this requirement remains a major open problem.
2. **Ideal lattice security.** While the best known algorithms for SVP on ideal lattices to within polynomial factors require time $2^{\Omega(\sqrt{n})}$, it is unclear whether ideal lattices are fundamentally easier than general lattices. Understanding this gap is important for calibrating Ring-LWE parameters.

7 Things I Learned

- I initially found it unclear why Ring-LWE requires passing to the number field rather than working purely in the ring. I learned that the relevant geometry does not live in the ring itself, but in its canonical embedding.
- In particular, the error distribution is naturally defined via the dual lattice, corresponding to the codifferent ideal $R^\vee$, which requires working in the embedded space to obtain the correct scaling and Fourier-analytic behavior.
- This clarified how number fields provide a rich enough setting to define lattice problems, with embeddings enabling a geometric interpretation necessary for worst-case reductions.
- I also learned how algebraic notions such as orders and trace maps allow us to generalize these constructions and move between related problems in a structured way.
- Finally, I saw that the additional algebraic structure in Ring-LWE introduces subtle challenges: reductions become more complicated, and establishing properties like search-to-decision equivalence requires more careful arguments than in standard LWE.

8 Updates

- Peikert, Regev, and Stephens-Davidowitz [5] give a direct quantum reduction from worst-case ideal lattice problems to *decision* Ring-LWE, extending hardness to arbitrary number fields and moduli and eliminating the need for search-to-decision reductions.

9 Conclusion

Starting from Regev’s foundational quantum reduction tying LWE hardness to worst-case lattice problems, we traced the development of Ring-LWE (which achieves compact keys via polynomial ring structure), and the further generalizations to Module-LWE, Polynomial-LWE, and Order-LWE. The unifying framework of “Algebraic LWE Revisited” shows that Ring-LWE reduces to these variants, providing a clean and general hardness foundation. These problems underpin the NIST post-quantum cryptography standards now being deployed worldwide, as well as modern FHE schemes, underscoring both their theoretical depth and practical importance. Algebraic variants provide more structure and can increase efficiency, making these protocols practical for real-world use.

References

- [1] O. Regev, *On Lattices, Learning with Errors, Random Linear Codes, and Cryptography*, Journal of the ACM, 56(6), 2009. (Preliminary version in STOC 2005.)
- [2] V. Lyubashevsky, C. Peikert, and O. Regev, *On Ideal Lattices and Learning with Errors Over Rings*, Journal of the ACM, 60(6), 2013. (Preliminary version in EUROCRYPT 2010.)
- [3] C. Peikert, *How (Not) to Instantiate Ring-LWE*, Cryptology ePrint Archive, Report 2016/504, 2016.
- [4] C. Peikert and Z. Pepin, *Algebraically Structured LWE, Revisited*, Journal of Cryptology, 37, 28, 2024. (Preliminary version in TCC 2019.)
- [5] C. Peikert, O. Regev, and N. Stephens-Davidowitz, *Pseudorandomness of Ring-LWE for Any Ring and Modulus*, Proceedings of the 49th Annual ACM SIGACT Symposium on Theory of Computing (STOC), 2017.